

# Assignment on Neural networks and mixture of Gaussian

Anass belcaid

2019/01/26

## Contents

<i>Backpropagation for Convnets</i>	1
<i>Neural Networks</i>	2
<i>Basic Generalization</i>	2
<i>Classification error</i>	3
<i>Learning rate</i>	4
<i>Number of hidden units</i>	4
<i>Compare k-NN and Neural Network</i>	4

Overview In this assignment, you will experiment with a neural network and the mixture of *Gaussians* model. Some code that implements a neural network with one hidden layer, and the mixture of Gaussians model will be provided for you (both MATLAB and Python). You will be working with the following dataset: Digits: The file `digits.mat` contains 6 sets of  $16 \times 16$  greyscale images in vector format (the `jpixel` intensities are between 0 and 1 and were read into the vectors in a raster-scan manner). The images contain centered, handwritten 2's and 3's<sup>1</sup>, scanned from postal envelopes. `train2` and `train3` contain examples of 2's and 3's respectively to be used for training. There are 300 examples of each digit, stored as  $256 \times 300$  matrices. Note that each data vector is a column of the matrix. `Valid2` and `valid3` contain data to be used for validation (100 examples of each digit) and `test2` and `test3` contain test data to be used for final evaluation only (200)

<sup>1</sup> This turns out to be a **binary** classification problem instead of  $K$  classes.

## Backpropagation for Convnets

You are training a Convolutional Neural Network (CNN) by minimizing the cross-entropy:

- The input is  $32 \times 32$  image<sup>2</sup>
- The first (and only) hidden layer is **convolutional**. There are  $F$  number of filters with size  $w \times h$ . The activation function is **RELU**.
- The output layer is *fully-connected* and has 3 units. It has the **soft-max** activation function.

<sup>2</sup> grayscale, not RGB

How many weights are there in the model? Explain how back propagation works, and derive equations for the updates for each

weight in the model. How many operations does the forward pass require?

### *Neural Networks*

Code for training a neural network with one hidden layer of logistic units, logistic output units and a cross entropy error function is included.

- **nn.py** : contains all the methods for initializing, training and validating the model.

### *Basic Generalization*

Train a neural network with 10 hidden units. You should first use `Init` to initialize the net, and then execute `train nn` repeatedly (more than 5 times). Note that `train nn` runs 100 epochs each time and will output the statistics and plot the error curves. Alternatively, if you wish to use Python, set the appropriate number of epochs in `nn.py` and run it. Examine the statistics and plots of training error and validation error (generalization). How does the network's performance differ on the training set versus the validation set during learning?

**Both errors are decreasing and (as expected) the training error is always inferior the validation and test error.**

Show a plot of error curves (training and validation) to support your argument.




Figure 1: img

Caption : Cross entropy loss for a simple Neural network with a hidden layer

### Classification error

You should implement an *alternative performance* measure to the cross entropy, the **mean classification error**<sup>3</sup>. You should then count up the total number of examples that are classified incorrectly according to this criterion for training and validation respectively, and maintain this statistic at the end of each epoch. Plot the classification error vs.

<sup>3</sup> You can consider the output correct if the correct label is given a higher probability than the incorrect label

Here is the code of the added function to compute the classification rate:


```
def classification_rate(inputs, target, W1, W2, b1, b2):
    """ Evaluat the classification rate on the model"""

    h_input = np.dot(W1.T, inputs) + b1 # Input to hidden layer.
    h_output = 1 / (1 + np.exp(-h_input)) # Output of hidden layer.
    logit = np.dot(W2.T, h_output) + b2 # Input to output layer.
    prediction = 1 / (1 + np.exp(-logit)) # Output prediction.

    #max prediction
    prediction = (prediction <= 0.5) #if proability higher than 0.5 class 1

    return 1 - np.mean(prediction != target)
```

Figure 2: img2



caption: Classification error rate.

*Learning rate*

Try different values of the **learning rate**  $\epsilon$ .

- You should reduce it to 0.01.
- Increase it to 0.2 and 0.5.

What happens to the convergence properties of the algorithm?

- Also try different values of the momentum in  $\{0.0, 0.5, 0.9\}$ .

**Generally we use cross-validation to set the best parameter.**

*Number of hidden units*

Set the learning rate  $\epsilon = 0.02$ , momentum to 0.5 and try different number of hidden units on the problem <sup>4</sup>.

Describe the effect of this modification on the convergence properties, and the generalization of the network.

With this setup, the model lacks the number of iteration to converge since the learning rate is **small** and the maximum allowed *epochs* is also reduced to 100.

<sup>4</sup> You should use two values  $\{2, 5\}$  which are smaller than the original two others  $\{30, 100\}$

*Compare  $k$ -NN and Neural Network*

Try  $k$ -NN on this digit classification task using the code developed in the first assignment<sup>5</sup>, and compare the result with those you got using neural networks. Briefly comment on the differences between these classifiers.

<sup>5</sup> I'm intrigued to use *scikit-learn*.

```
inputs_train, inputs_valid, inputs_test, \
target_train, target_valid, target_test \
= LoadData('digits.npz')

knn_valid_target = run_knn(K, inputs_train.T \
, target_train.T, inputs_valid.T).squeeze()
knn_test_target = run_knn(K, inputs_train.T \
, target_train.T, inputs_test.T).squeeze()

knn_valid_error = 1 - mean(knn_valid_target == target_valid)
knn_test_error = 1 - np.mean(knn_test_target == target_test)

print("{:2d}-NN validation error={:4.2f}, test error={:4.2f}" \
.format(K, knn_valid_error, knn_test_error))
```

We obtain a similar error 2% with  $k = 10$ . This error increases with a higher  $k$ .